

Assignment 1: Command-line Memory

CSE 130: Principles of Computer Systems Design

Due: October 15, 2024 at 11:30 PM

Goal This project serves as a refresher for building software in ‘C’. In particular, the project will involve a review of Linux system calls, buffering file I/O, memory management, and C-string parsing.

Assignment Details

For this assignment, you will write a program, `memory`, that provides a `get/set` memory abstraction for files in a Linux directory. Your program should take a command from `stdin` and carry out the command in the current working directory. Below, we provide examples of how `memory` should work, outline the commands formally, describe how to handle invalid commands, and provide notes about other functionality and limitations of `memory`.

Examples

Here, we illustrate how your server should operate with four examples. In each example, suppose that the user starts `memory` from a directory that includes the file `foo.txt`, with the content “Hello from foo”, and the file `bar.txt`, with the content “Hello from bar”.

1. Example 1. In this example, the user passes the string “`get\nfoo.txt\n`” to `stdin`. `memory` outputs the contents of `foo.txt` (“Hello from foo”) to `stdout` and then exits with a return code of 0.
2. Example 2. In this example, the user passes the string “`set\nbaz.txt\n12\nHello World!`” to `stdin`. `memory` should create a new file, named `baz.txt`, in the current directory, write “Hello World!” to `baz.txt`, write the message OK to `stdout`, and then exit with a return code of 0.
3. Example 3. In this example, the user passes the string “`set\nfoo.txt\n12\nHello World!`” to `stdin`. `memory` should change the contents of `foo` to “Hello World!”, write the message OK to `stdout`, and then exit with a return code of 0.
4. Example 4. In this example, a user passes the string “`invalid\nfoo.txt`” to `stdin`. `memory` should output the message “Invalid Command\n” to `stderr`, and exit with a return code of 1.

Commands

Your program will receive a command from `stdin` as input. Each valid command will be either a `get` or `set` command.

- `get`. A valid `get` command is formatted as “`get\n<location>\n`” and indicates that `memory` should output the contents of the file named `location`. If there is a file named `location` in the current directory, then `memory` should write the contents of the `location` to `stdout` and exit with a return code of 0.
- `set`. A valid `set` command is formatted as “`set\n<location>\n<content_length>\n<contents>`” and indicates that `memory` should assign `content_length` bytes from `contents` to a file named `location` in the current directory. If the `location` is a valid filename (see Hints section for some tips on how to check for this), then `memory` should assign the content of the `location` to `content_length` bytes from `content` (overwriting the `location`’s current contents if applicable) and exit with return code 0. In general, `memory` should be permissive on input—the Notes section gives details on how to interpret this.

Invalid Commands

If `memory` receives an invalid command, then it should output the text “Invalid Command\n” to `stderr` and exit with a return code of 1. An invalid command meets any of the following:

- The command does not start with either `get` or `set` (these are case-sensitive).
- The command starts with `get` or `set`, but `stdin` was closed before the user provided a `location`.
- The command starts with `get` or `set`, and includes a `location`, but the `location` is not a valid filename (see Hints section at the end of this document).
- The command is a `get` command, and includes a valid `location`, but there is no file located at the `location`.
- The command is a `get` command, includes a valid `location` that exists, but also includes additional input (e.g., “`get\nfoo.txt\nExtraStuff`”).

Notes

In addition to the functionality described above, your `memory` should meet the following functionality and limitations.

- If `memory` detects any other error (e.g. if it cannot write all of the requested content to a file for any reason), the program should produce the text `Operation Failed` to `stderr` and exit with a return code of 1.
- `memory` must be reasonably time efficient: it should buffer file input and output.
- `memory` must be reasonably space efficient: it should use at most 10 MB of memory regardless of input.
- `memory` must be permissive on input: if a user issues a `set` command and closes `stdin` before providing `content_length` bytes, then the file at the `location` should contain exactly the `content` provided in the command (i.e., `memory` should allow a user to provide too few bytes). Similarly, if a user issues a `set` command and includes more than `content_length` bytes, then the file at the `location` should contain exactly the first `content_length` bytes from `content` provided in the command (i.e., `memory` should allow a user to provide too many bytes).
- `memory` must not leak any memory (i.e., it should free all memory that it allocates).
- `memory` must not leak any file descriptors (i.e., it should close all files that it opens).
- `memory` should not crash (e.g., it should never `segfault`).
- `memory` must be written using the ‘C’ programming language (*not C++!*).
- `memory` cannot use the following functions from the ‘C’ `stdio.h` library: `fwrite`, `fread`, variants of `put` (i.e., `fputc`, `putc`, `putc_unlocked`, `putchar`, `putchar_unlocked`, and `putw`), and `get` (i.e., `fgetc`, `getc`, `getc_unlocked`, `getchar`, `getchar_unlocked`, `getline`, and `getw`).
- `memory` cannot use functions, like `system(3)`, that execute external programs.

Rubric

We will use the following rubric for this assignment:

Category	Point Value
Makefile	10
Clang-Format	5
Files	5
Functionality	80
Total	100

Makefile Your repository includes a Makefile with the rules `all` and `memory`, which produce the `memory` binary, and the rule `clean`, which removes all `.o` and binary files. Additionally, your Makefile should use `clang` (i.e., it should set `CC=clang`), and should use the `-Wall`, `-Wextra`, `-Werror`, and `-pedantic` flags (i.e., it should set `CFLAGS=-Wall -Wextra -Werror -pedantic`).

Clang-Format All `.c` and `.h` files in your repository are formatted by the `.clang-format` file included in your repository.

Files The following files are included in your repository: `memory.c`, `Makefile`, and `README.md`. Your repository also should not include binary files or any object files (i.e., `.o` files). To make it easier for you to maintain tests, we added an exception to the binary file rule: you *can* include binary files in any directory whose name starts with the phrase `test`.

Functionality Your `memory` program performs the functionality described in the Assignment Details section.

How to submit

1. Navigate to the `asgn1` directory in your repo.
2. Configure your `config.json` file (see Canvas or the README in the resources repo for more information).
3. `python3 -m autograder.run.submit [file1] [file2] ...`

Resources

Here are some resources to help you:

Testing We provided you with two resources to test your own code:

1. An autograder, which you will run locally through the command line.
2. A set of test scripts in the resources repository to check your functionality. You can use the tests to see if your functionality is correct by running them on your Ubuntu 24.04 virtual machine. We provided you with a subset of the tests that we will run, but, I bet you can figure the other ones out by adapting what we have given you :-)

Hints

- Check out the document on Canvas about setting up Ubuntu. It has instructions for getting an Ubuntu 24.04 virtual machine up and running using Multipass.
- To format a file, named `foo.c`, using the included `.clang-format`, execute the command `clang-format -i -style=file foo.c`. If you include the `.clang-format` file in the root directory of your repository (as it was when we created your repositories), you can instead execute `clang-format -i foo.c`.
- System calls that allow you to interact with standard input and output (e.g. `read` and `write`) will be helpful to achieve the functionality listed above. If you aren't sure how to use them, look up a function's manual page by typing `man <command>` in your terminal. Or, use google!
- To check the return code of your program, run your program until completion and then type `echo $?` in your terminal.
- You can check your program for memory leaks using `valgrind --leak-check=full`.

- `memory` should allow a user to `get` or `set` data from any filename that is valid on a Linux filesystem. In Linux, the only restrictions on filenames are that they cannot be longer than `PATH_MAX` from the "`<limits.h>`" header and that they cannot include null characters (i.e., `\0`). To make your life easier, we're simplifying further: none of our tests send input to your `memory` program and specify a valid method (i.e., either `get` or `set`) will include a filename with non-ASCII characters or spaces.
- This assignment requires that you perform some basic string parsing. There are many ways to parse a C-string; the most straightforward for this assignment is probably `strtok`, but you could also check out `regex` or `sscanf`.